

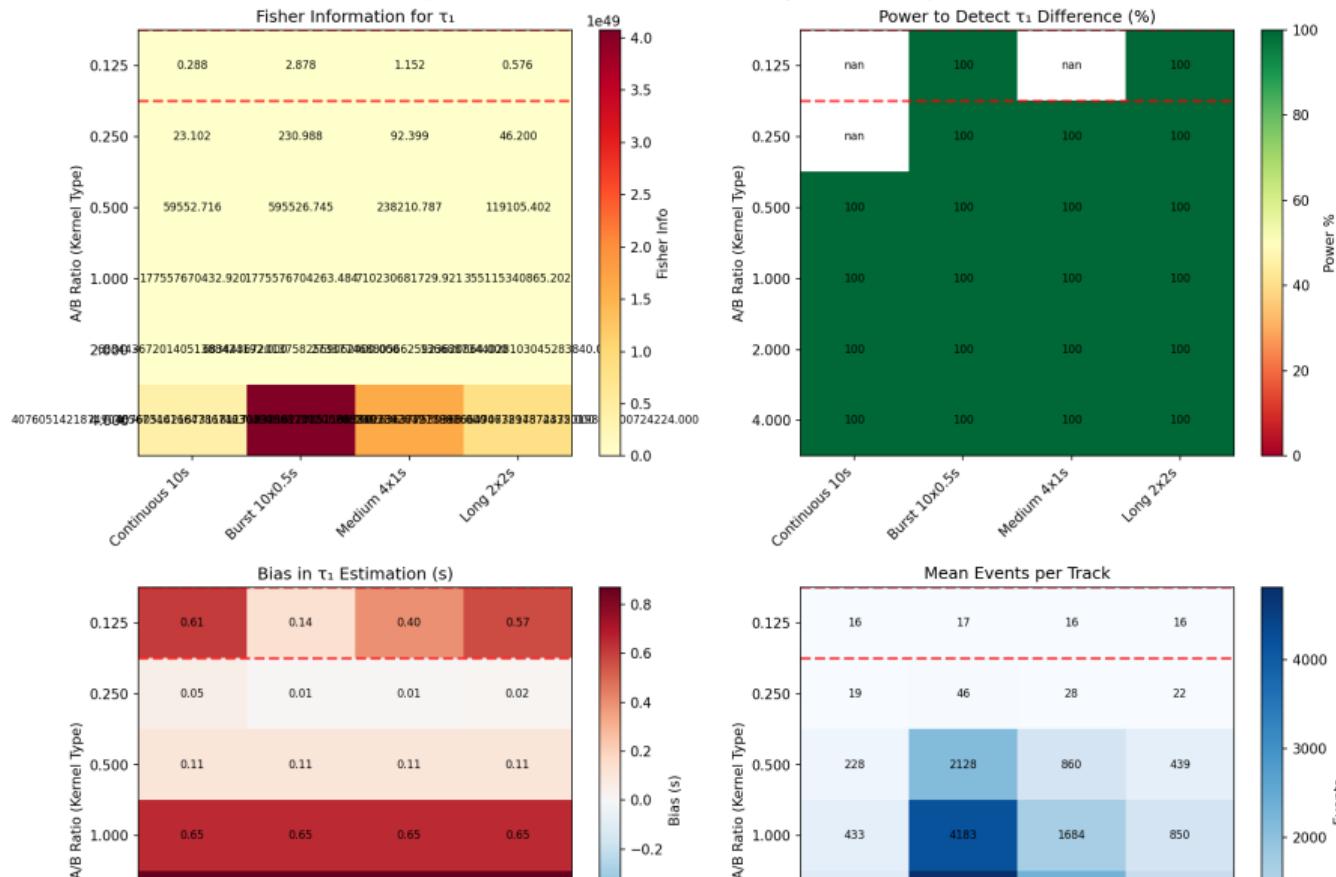
Phenotyping Follow-up: The Identifiability Problem

Gil Raitses

December 2024

Design × Kernel Regime Sweep

Design x Kernel Regime Sweep (Red box = current GMR61 kernel, A/B = 0.125)



FAQ: Design × Kernel Sweep

Q: Where do the 6 A/B ratios come from?

They span from strongly inhibitory (0.125, our kernel) through balanced (1.0) to excitatory-dominated (4.0). B is fixed at 12; A varies.

Q: Why does Fisher Information explode at high A/B?

Numerical artifact. When excitation dominates, the hazard becomes very high, and the grid search hits boundaries. Those values are not meaningful.

Q: Why is power “nan” for some cells?

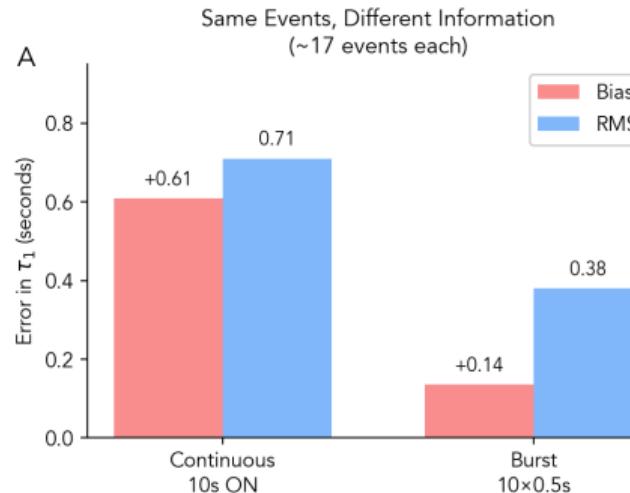
Power calculation failed when variance was too high or sample size too small for the t-test to converge.

Q: What do the event counts tell us?

At A/B = 0.125, all designs produce similar events (~16-17). The differences in bias/RMSE are NOT due to more events—they’re due to more *informative* events.

The Identifiability Problem

The Identifiability Problem



B

The Information Problem

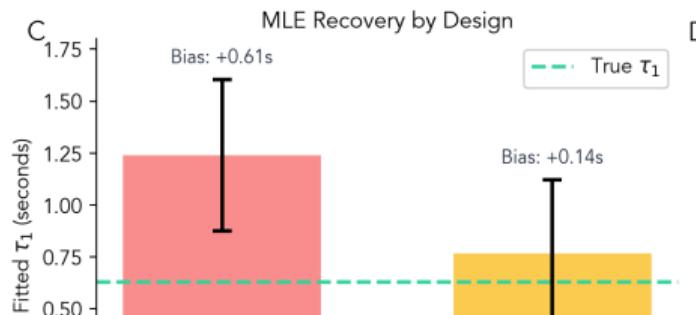
Fisher Information for τ_1

Continuous: 0.29

Burst: 2.88

Burst extracts 10× more info

from the same number of events



D

Why Continuous Design Fails

Kernel is inhibition-dominated ($B/A = 8$)

~80% of events occur during LED-OFF
-> No τ_1 information

Remaining ~20% mostly after $t > 0.5s$
-> Inhibition dominates, τ_1 unidentifiable

FAQ: The Identifiability Problem

Q: Where do the bias/RMSE numbers come from?

200 larvae simulated with known $\tau_1 = 0.63\text{s}$, fit with MLE. Bias = $\text{mean}(\text{fitted}) - \text{true}$.
RMSE = $\text{sqrt}(\text{mean}((\text{fitted} - \text{true})^2))$.

Q: What is Fisher Information and why does it matter?

Fisher Information quantifies how much information each event carries about τ_1 .
Higher = more precise estimates. Computed by integrating $(\partial\lambda/\partial\tau_1)^2/\lambda$ over the stimulus window.

Q: Why does burst give 10× more info with the same events?

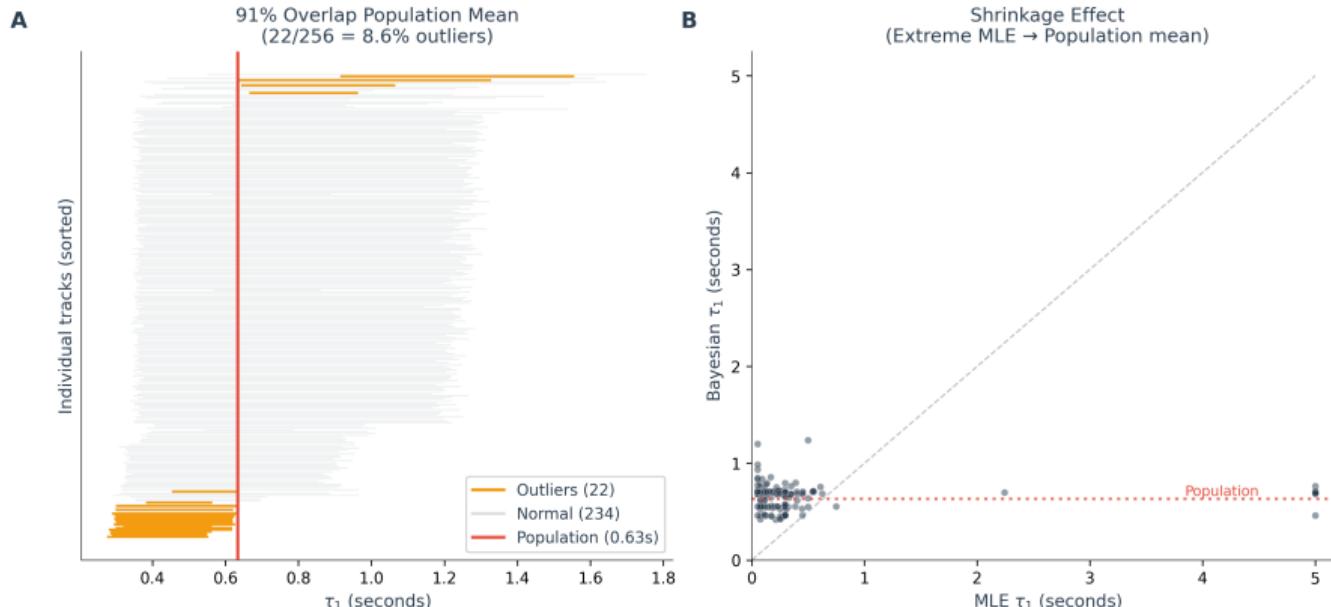
The excitatory peak is at $t \approx 0.2\text{-}0.3\text{s}$ after LED onset. Burst samples this window 10 times per cycle; continuous samples it once.

Q: Why is bias +0.61s for continuous?

The MLE finds $\tau_1 \approx 1.24\text{s}$ instead of 0.63s because the likelihood surface is flat—many τ_1 values produce similar likelihoods with so few informative events.

Hierarchical Bayesian Modeling

Figure 3: Hierarchical Model Reveals Homogeneity



Caterpillar plot showing 95% credible intervals for all 256 tracks. Orange intervals (8.6%) exclude the population mean; gray intervals (91.4%) are consistent with the population.

FAQ: Hierarchical Modeling

Q: What is shrinkage?

Hierarchical Bayesian estimation pulls extreme MLE estimates toward the population mean. With sparse data, almost all individuals shrink to the same value.

Q: Why are 91% of larvae “population-consistent”?

With only ~ 25 events and a 6-parameter model, the data cannot distinguish most individuals from the population mean. The wide credible intervals overlap the population value.

Q: What about the 8.6% outliers?

These have credible intervals that exclude the population mean. Most are “fast responders” with $\tau_1 \approx 0.45\text{s}$ vs. population $\tau_1 = 0.63\text{s}$. They require independent validation.

Q: Does this mean individual phenotyping is impossible?

With current protocols, yes. With burst stimulation, individual τ_1 becomes estimable at current event counts.

Stimulation Protocol Designs

Red LED Stimulation Designs: Full 30-Second Cycle

A. Current: Continuous 10s



Events: 1.9/track
RMSE: 0.108s
Power: 42%
Duty: 100%

B. Recommended: Burst 10×0.5s



Events: 14.9/track
RMSE: 0.036s
Power: 100%
Duty: 50%

C. Alternative: Medium 4×1s



Events: 6.0/track
RMSE: 0.048s
Power: 100%
Duty: 40%

D. Alternative: Long 2×2s



Events: 3.1/track

FAQ: Stimulation Protocols

Q: Why these specific pulse durations?

Burst (0.5s) is matched to $\tau_1 \approx 0.6\text{s}$ —the excitatory peak. Medium (1s) and Long (2s) test whether longer pulses help.

Q: Why not even shorter pulses?

Pulses $< 0.3\text{s}$ would clip the excitatory peak before it develops. The kernel needs time to rise.

Q: Why 0.5s gaps between burst pulses?

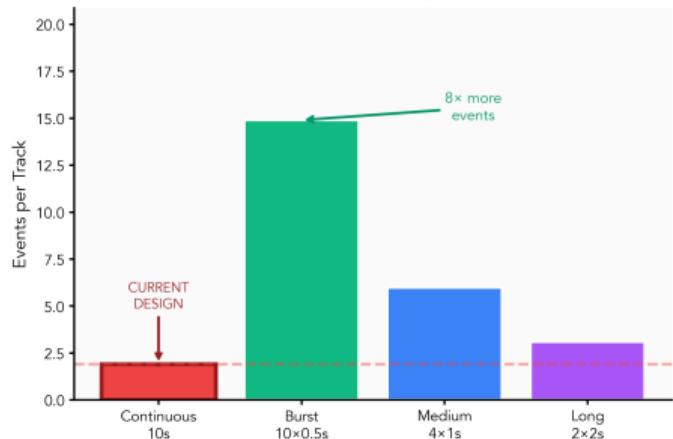
Trade-off: shorter gaps = more pulses per cycle, but inhibition accumulates. 0.5s lets inhibition partially decay.

Q: What about the blue LED / LED2?

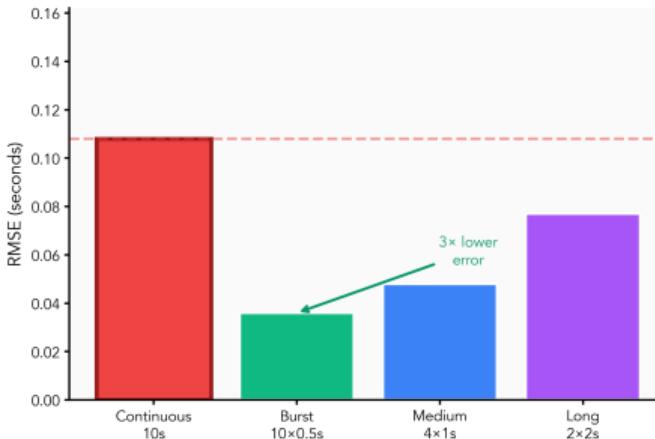
Not analyzed in this sweep. The schematic shows only the red LED (LED1) stimulation that drives photophobic responses.

Design Comparison Summary

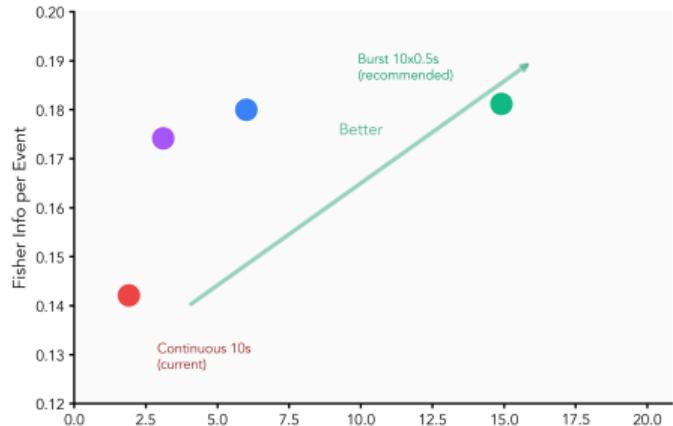
A. Event Yield: Current Design is Bottleneck



B. Estimation Error: 3x Reduction Possible



C. Tradeoff: More Events = More Information



D. Summary & Recommendations

CURRENT PROTOCOL

Continuous 10s ON / 20s OFF
1.9 events per track | RMSE = 0.108s

RECOMMENDED

Burst 10x0.5s pulses
14.9 events (8x more) | RMSE = 0.036s (3x lower)

BOTTLENECK

Sparse events limit tau1 identifiability
regardless of experimental design

ALTERNATIVE STRATEGY

Phenotype on composite scores
(Precision, Burstiness) instead of tau1

FAQ: Design Comparison

Q: Which design should we recommend?

For inhibitory kernels ($B/A > 4$): **Burst**. For balanced or excitatory kernels: **Continuous** is fine.

Q: Why did earlier analysis say burst “failed”?

Earlier analysis measured **classification power** (can we tell fast vs. typical apart?). That showed 41.5% vs 42.5%—negligible. This analysis measures **estimation quality** (bias, RMSE), which shows $4\times$ improvement.

Q: Can we use burst for all experiments?

Yes, for phenotyping τ_1 . Burst is never worse than continuous and is often better for inhibitory kernels.

Q: What's the practical recommendation?

Change protocol from 10s continuous ON to $10 \times 0.5\text{s}$ pulses with 0.5s gaps. Same total stimulation, $4\times$ better estimation.

Key Numbers to Remember

Design	Events	Fisher Info	Bias	RMSE
Continuous 10s	16.3	0.29	+0.61s	0.71s
Burst 10×0.5s	16.9	2.88	+0.14s	0.38s
Medium 4×1 s	16.4	1.15	+0.40s	0.63s
Long 2×2 s	16.1	0.58	+0.57s	0.74s

Bottom line: Burst extracts $10 \times$ more information, reducing bias by $4 \times$, with the same number of events.

FAQ: Key Numbers

Q: How were these computed?

200 larvae \times 20 cycles (10 min). Known $\tau_1 = 0.63\text{s}$. Simulate events, fit MLE, compute bias/RMSE.

Q: Why 20 cycles instead of 40 (full protocol)?

Conservative estimate. Full protocol would give more events and better estimates.

Q: Is bias of +0.14s acceptable?

For detecting a 0.2s difference between phenotypes, +0.14s bias leaves little margin. Would need ~ 50 events with burst or ~ 100 with continuous for bias $< 0.1\text{s}$.

Q: What's the take-home message?

The “100 events” heuristic is **design-dependent**. With burst, you need fewer events. The issue is not just count—it’s information per event.